

CaRS: Causal Regime-Switching State Space Model for Electricity Price Forecasting

Combining Deep State Space Models with Causal Structure Learning

January 2026

- 1 Motivation
- 2 Technical Challenges
- 3 Economic Relevance
- 4 Data
- 5 Literature Review
- 6 Methodology
- 7 Numerical Results
 - Shared Backbone Experiments
 - Shared Backbone Estimation Experiments
- 8 Discussion
- 9 Outlook

Motivation: Understanding *causal* price drivers enables better policy decisions, more efficient markets, and reduced consumer costs.

Energy Security & Grid Stability

- Accurate price forecasts essential for **grid balancing**
- Misaligned supply/demand → blackouts or price spikes
- Renewable integration requires understanding price volatility

Market Efficiency & Price Discovery

- Transparent pricing mechanisms
- Efficient allocation of generation resources
- Cross-border trading optimization (DE-FR interconnection)

Financial Risk Management

- Utilities hedge against price volatility
- Trading firms require reliable forecasts
- Investment decisions for new generation capacity

Policy & Consumer Protection

- Carbon pricing impact assessment
- Renewable subsidy effectiveness
- Consumer electricity bill forecasting
- Regulatory compliance (REMIT, EU energy law)

Why Electricity Price Estimation & Prediction?

Estimation Task:

- Given: $X_{0:t}, Y_{0:t-1}$
- Estimate: Y_t (today's price)
- **Use case:** Real-time market analysis, missing data imputation

Prediction Task:

- Given: $X_{0:t}, Y_{0:t}$
- Predict: Y_{t+1} (tomorrow's price)
- **Use case:** Trading decisions, risk management

Why Switching Regimes?

- Energy markets exhibit **structural breaks**
- Different price dynamics during:
 - Peak vs. off-peak hours
 - Summer vs. winter
 - Normal vs. crisis periods
- **Causal relationships change** across regimes
- Single-regime models miss important dynamics

Goal: Learn regime-specific causal structures that explain *which factors drive prices and how this changes over time*

1. Latent/Unobserved Variables

- Market regimes are **not directly observable**
- Must infer regime from price behavior
- Regime transitions follow **Markov dynamics**

2. Non-Stationary Transition Dynamics

- Regime transition probabilities change
- $P(d_t | d_{t-1})$ varies with market conditions
- Standard Hidden Markov Models (HMMs) assume fixed transitions

3. High-Frequency Data

- Hourly electricity prices
- Strong **autocorrelation** structures
- Intraday and weekly seasonality

4. Non-Linear Dependencies

- Complex interactions between:
 - Renewable generation (wind, solar)
 - Commodity prices (gas, coal, carbon)
 - Load/demand patterns
- Linear models insufficient

Challenge: How to jointly learn (1) latent regimes, (2) regime-specific causal graphs, and (3) non-linear dynamics?

CaRS addresses common challenges across multiple domains:

Financial Markets

- Regime shifts: bull/bear markets, volatility clusters
- Non-stationary dependencies between assets
- Causal drivers change during crises

Healthcare & Epidemiology

- Patient monitoring with changing conditions
- Disease progression with treatment regimes
- Epidemic dynamics (normal vs. outbreak)

Climate Science

- Weather pattern regime shifts (El Niño/La Niña)
- Non-linear climate feedback mechanisms
- Changing causal relationships over seasons

Manufacturing & IoT

- Process monitoring with multiple operating modes
- Quality control under varying conditions
- Predictive maintenance with regime detection

Why Causal Models for Decision Making?

Interpretability

- Understand *why* prices change
- Identify key drivers per regime
- Communicate to stakeholders
- Regulatory compliance


Trustworthiness

- Causal graphs can be validated by domain experts
- Predictions grounded in economic theory
- Detect spurious correlations
- Build confidence in model outputs

Robustness

- Causal models generalize better under **distribution shift**
- Regime-specific DAGs adapt to changing conditions
- More stable predictions during market stress

Why Causal Models for Decision Making?



Benefit for Energy Trading: Knowing that “wind generation $\rightarrow$ price” is causal (not just correlated) enables better hedging strategies and scenario analysis under different weather conditions.

Target Variables:

- **DE:** German day-ahead price (EUR/MWh)
- **FR:** French day-ahead price (EUR/MWh)
- **DE_FR:** Price spread

Explanatory Features (19 total):

- **Generation (9):** Wind, Solar, Hydro, Nuclear, Gas, Coal, Lignite, Biomass (MW)
- **Load (1):** Actual demand (MW)
- **Commodities (3):** Natural Gas, Brent, WTI (prices)
- **Weather (6):** Temperature, Wind, Radiation, Cloud, Precipitation

Data Statistics (Daily, 2015-2024):

	DE	FR	Spread
Samples	3,652	3,652	3,651
Mean	71 EUR	78 EUR	-7 EUR
Std	78 EUR	88 EUR	27 EUR
Min	-50 EUR	-9 EUR	-480 EUR
Max	705 EUR	740 EUR	222 EUR

Key characteristics:

- Heavy tails (kurtosis: 13 DE, 12 FR, 42 spread)
- Strong autocorrelation (lag-1: 0.94, 0.97, 0.71)
- Prices right-skewed; spread left-skewed

Power Generation Mix: Energy Transition (2015 → 2024)

Germany (DE)

2015

2024

France (FR)

2015

2024

Key changes: DE: Nuclear phase-out (14% → 0%), wind doubles (21% → 41%). FR: Nuclear stable (~70%).

Time Series: Germany Day-Ahead Prices (2015-2024)

Time Series: France Day-Ahead Prices (2015-2024)

Model Comparison: Capabilities

Model	Regime Switching	Causal DAGs	Temporal Dynamics	Non-Linear Relations	Interpretable Graphs	Scalable (>50 feat.)
ARIMA/VAR	✗	✗	✓	✗	✗	✓
MS-VAR	✓	✗	✓	✗	✗	✓
NOTEARS/GOLEM	✗	✓	✗	✓	✓	✓
DYNOTEARS	✗	✓	✓	✗	✓	✓
DS3M	✓	✗	✓	✓	✗	✓
FANTOM	✓	✓	✓	✓	✓	✗
CaRS	✓	✓	✓	✓	✓	✓

Key References:

- **DS3M** (Xu et al., 2021): Deep Switching State Space Model for time series
- **FANTOM** (Rahmani et al., 2025): Causal discovery with regime switching
- **NOTEARS** (Zheng et al., 2018): Differentiable DAG learning

From DS3M (Temporal Backbone):

- GRU encoder: $h_t = \text{GRU}(h_{t-1}, [X_t, Y_t])$
- Learnable Markov regime switching
- Continuous latent state z_t

From FANTOM (Causal Discovery):

- ICGNN emission network
- Gumbel-Softmax DAG sampling
- NOTEARS acyclicity + sparsity constraints

CaRS Novelty — Joint End-to-End Learning:

The GRU hidden state h_t simultaneously conditions:

- 1 the **regime posterior** $P(d_t | h_t, d_{t-1})$ — *which* causal structure is active
- 2 the **causal emission** $\hat{Y}_t = f_{\text{ICGNN}}(X_t, h_t, z_t, A^{(d_t)})$ — *how* to predict through the DAG

Regimes, causal DAGs, and predictions are learned in a **single ELBO optimization** — not a two-stage pipeline.

DS3M_UV (Univariate)

- Only uses target Y
- Autoregressive regime switching
- Strong baseline for price forecasting

DS3M_MV (Multivariate)

- Uses features X and target Y
- Black-box emission network
- No interpretable structure

CaRS (Ours)

- Combines SSM regime switching with causal discovery
- **Key innovation:** Enhances SSM with Invertible Contractive Graph Neural Network (ICGNN)
- Learns **sparse causal DAGs per regime**

FANTOM_BEM (Baseline)

- Bayesian EM for regime detection
- Full FANTOM causal discovery
- Reference implementation

Novelty: CaRS jointly learns *regime-switching causal DAGs* end-to-end within a deep state space—regimes and causal structure are discovered simultaneously, not in a two-stage pipeline. Each regime defines a Structural Equation Model (ICGNN + NOTEARS), so every prediction decomposes into *interpretable directed causal edges* that change with market regime.

Key Methodological Differences: DS3M vs FANTOM vs CaRS

Component	DS3M	FANTOM	CaRS (Ours)
Temporal Encoding	GRU	Lag-based (no recurrence)	GRU
Regime Detection	Markov (learnable)	Bayesian EM	Markov (learnable)
State Space	Continuous z_t	None	Continuous z_t
Emission Model	Black-box MLP	ICGNN	ICGNN
Causal DAG Learning	None	Gumbel-Softmax (per-regime)	Gumbel-Softmax (per-regime)
Noise Modeling	Gaussian	Cond. Norm. Flows	Gaussian*

*CaRS uses Gaussian noise for computational efficiency and numerical stability, enabling scalable training.

CaRS Innovation:

From DS3M:

- GRU temporal encoding
- Learnable Markov regimes
- Continuous latent state


From FANTOM:

- ICGNN architecture
- Differentiable DAG learning
- Sparsity constraints

- h_t conditions regime posterior AND causal emission
- Joint end-to-end ELBO optimization
- Temporal + causal +


The Joint Learning Mechanism

CaRS: Joint End-to-End



h_t informs **both** which regime is active **and** how to predict through the causal graph.

FANTOM: Two-Stage Pipeline



Regimes detected from Y alone; DAGs learned independently per regime.

Key difference: In CaRS, temporal context (h_t) directly shapes causal structure — the model discovers that causal relationships change *because of* temporal dynamics, not just segments data into static regimes.

Observed Variables:

- $X_t \in \mathbb{R}^n$: **Exogenous features** at time t
 - Generation: Wind, Solar, Gas, Nuclear (MW)
 - Commodities: Gas price, Brent, Carbon (EUR)
 - Weather: Temperature, Wind speed, Radiation
- $Y_t \in \mathbb{R}$: **Target variable** (electricity price, EUR/MWh)

Data Structure:

- Time series:
 $(X_1, Y_1), (X_2, Y_2), \dots, (X_T, Y_T)$
- Dimensions: $n = 19$ features, $T \approx 3,652$ days

Tasks:

- 1 **Estimation:** Given $X_{0:t}, Y_{0:t-1}$, estimate Y_t
- 2 **Prediction:** Given $X_{0:t}, Y_{0:t}$, predict Y_{t+1}
- 3 **Causal Discovery:** Which $X_t^{(i)} \rightarrow Y_t$?

Key Challenge:

Causal relationships *change over time!*

E.g., Wind's impact on price differs
in normal vs. crisis market conditions.

Y_t is generated by a **regime-specific** structural equation:

$$Y_t = f^{(d_t)}(\text{Pa}_{G^{(d_t)}}(Y_t)) + g^{(d_t)}(\text{Pa}_{G^{(d_t)}}(Y_t)) \cdot \epsilon_t$$

Step-by-step breakdown:

- 1 $d_t \in \{1, \dots, K\}$: Current **regime**
- 2 $G^{(d_t)}$: The **causal DAG** active in regime d_t
- 3 $\text{Pa}_{G^{(d_t)}}(Y_t) \subseteq X_t$: **Causal parents** of Y_t in the DAG
- 4 $f^{(d_t)}(\cdot)$: **Mean function** (ICGNN)
- 5 $g^{(d_t)}(\cdot)$: **Variance function** — captures heteroscedasticity
- 6 $\epsilon_t \sim \mathcal{N}(0, 1)$: Standardized noise

Example: If $\text{Wind} \in \text{Pa}(Y_t)$ but $\text{Solar} \notin \text{Pa}(Y_t)$, then wind *causally* affects price while solar does not (enforced via DAG mask $A^{(d_t)}$).

Latent Variables (per timestep t):

$$h_t = \text{GRU}(h_{t-1}, [X_t, Y_t]) \quad (\text{hidden state})$$

$$d_t \in \{1, \dots, K\} \quad (\text{discrete regime})$$

$$z_t \sim \mathcal{N}(\mu(h_t, d_t), \sigma(h_t, d_t)) \quad (\text{continuous latent})$$

Structural Parameters (per regime r):

$$A^{(r)} \in [0, 1]^{n \times n} \quad (\text{causal adjacency})$$

$$\theta^{(r)} = \{\psi^r, \vartheta^r\} \quad (\text{ICGNN weights})$$

Transition Parameters:

$$\Phi \in \mathbb{R}^{K \times K} \quad (\text{transition logits})$$

$$P(d_t | d_{t-1}) = \text{softmax}(\Phi)_{d_{t-1},:}$$

Emission (Final Prediction):

$$\hat{Y}_t = f_{\text{ICGNN}}(X_t, h_t, z_t, A^{(d_t)})$$

Learnable: GRU weights, ICGNN $\theta^{(r)}$,
DAG params $\phi^{(r)}$, transitions Φ

The Optimization Challenge: Why Direct MLE is Impossible

Goal: Maximize the log-likelihood of observed data $(X, Y)_{1:T}$:

$$\log p_{\Theta}(Y_{1:T}|X_{1:T}) = \log \int \sum_{d_{1:T}} p(Y, d, z, G | X) dz$$

Problem: This integral is INTRACTABLE!

1. Latent Regimes d_t

- d_t is **unobserved**
- Must sum over K^T sequences
- E.g., $3^{3652} \approx 10^{1742}$

2. Continuous Latents z_t

- Must integrate over all z_t values
- Non-linear ICGNN: no closed form

3. Unknown DAGs $G^{(r)}$

- Structure is **also latent**
- 2^{n^2} possible graphs
- E.g., $2^{361} \approx 10^{109}$

Solution: Instead of maximizing $\log p(Y|X)$ directly,
maximize a tractable **lower bound** (ELBO)

Evidence Lower Bound (ELBO): Making Optimization Tractable

Key Identity (Jensen's Inequality):

$$\log p(Y|X) \geq \underbrace{\mathbb{E}_q[\log p(Y|X, z, d, G)]}_{\text{Reconstruction}} - \underbrace{\text{KL}[q(z, d, G) \parallel p(z, d, G)]}_{\text{Regularization}} = \text{ELBO}$$

The Variational Distribution q :

- $q(z_t|h_t, Y_t)$: **Gaussian** posterior for z_t
- $q(d_t|h_t, Y_t)$: **Categorical** posterior for regime
- $q_\phi(G^{(r)})$: **Factorized Bernoullis** for DAG edges

What's special about q ?

- Tractable: we can sample from it
- Learnable: parametrized by neural networks

Why “Lower Bound”?

- $\text{ELBO} \leq \log p(Y|X)$ **always**
- $\text{Gap} = \text{KL}[q \parallel p(\cdot|Y, X)] \geq 0$

When is the bound tight?

- When $q = p(z, d, G | Y, X)$ exactly
- In practice: optimize q to minimize gap

Training: Maximize ELBO w.r.t. model params Θ and variational params ϕ

CaRS ELBO (expanded):

$$\mathcal{L} = \underbrace{\mathbb{E}_q[\log p(Y|X, z, d, A)]}_{\text{Reconstruction}} - \underbrace{\text{KL}[q(z)||p(z)]}_{\text{Continuous}} - \underbrace{\text{KL}[q(d)||p(d|d_{t-1})]}_{\text{Discrete}} + \underbrace{\lambda_{\text{dag}} h(A) + \lambda_{\text{sp}} \|A\|_1}_{\text{DAG Constraints}}$$

Monte Carlo Estimation:

Reparametrization Trick (continuous z_t):

- 1 Sample $\epsilon \sim \mathcal{N}(0, I)$
- 2 Compute $z_t = \mu(h_t, d_t) + \sigma(h_t, d_t) \odot \epsilon$
- 3 Gradients flow through μ, σ !

Gumbel-Softmax (discrete DAG A):

- DAG edges are Bernoulli: $A_{ij} \in \{0, 1\}$
- Relaxation: $\text{softmax}\left(\frac{\log \pi + G}{\tau}\right)$
- $G \sim \text{Gumbel}(0, 1)$; $\tau = \text{temperature}$

$$\mathbb{E}_q[f(z)] \approx \frac{1}{N} \sum_{i=1}^N f(z^{(i)}), \quad z^{(i)} \sim q$$


- Training: $N = 1$ sample (unbiased gradient)
- Inference: $N > 1$ for uncertainty

DAG Acyclicity (NOTEARS):

$$h(A) = \text{tr}(e^{A \odot A}) - n = 0$$


Sampling and Reparametrization Tricks

DAG Acyclicity: Ensures no directed cycles in learned graph. See Appendix for detailed derivations of KL terms, Gumbel-Softmax, and NOTEARS constraint.




- Discrete regime $d_t \in \{1, \dots, K\}$: Markov transitions between market states
- Continuous latent $z_t \sim \mathcal{N}(\mu(h_t, d_t), \sigma(h_t, d_t))$: smooth within-regime dynamics
- Regime-specific DAG $A^{(d)}$: learned via Gumbel-Softmax
- ICGNN: causal prediction using X_t and structure $A^{(d)}$

CaRS: Detailed Architecture with Shared Backbone



CausalICGNN: Internal Architecture



Hidden State h_t (GRU output)

- Encodes **temporal memory** from past observations
- $h_t = \text{GRU}(h_{t-1}, [X_t, Y_t])$
- Captures autocorrelation and trends
- Deterministic given input sequence

Discrete Regime d_t (Markov)

- Indicates **which market regime** is active
- $d_t \in \{1, 2, \dots, K\}$
- Transitions: $P(d_t | d_{t-1}) = \text{softmax}(\Phi)$
- Selects which causal DAG $A^{(d)}$ to use

Continuous Latent z_t (Gaussian)

- Captures **smooth within-regime variation**
- $z_t \sim \mathcal{N}(\mu(h_t, d_t), \sigma(h_t, d_t))$
- Enables **Monte Carlo uncertainty quantification**
- At inference: sample n times $\rightarrow$ mean $\pm$ std

How They Interact:

- 1 h_t summarizes temporal context
- 2 d_t determines regime (DAG selection)
- 3 z_t adds within-regime variability
- 4 ICGNN combines all to predict $\hat{Y}_t$

$$\hat{Y}_t = f_{\text{ICGNN}}(X_t, h_t, z_t, A^{(d_t)}) \text{ where } A^{(d_t)} \text{ is the regime-specific causal graph}$$

Uncertainty Quantification:

- CaRS provides **explicit predictive uncertainty** via Gaussian output
- Training: $p(Y_t|\cdot) = \mathcal{N}(\mu_t, \sigma_t)$ with learned heteroscedastic variance
- Inference: Monte Carlo sampling over z_t yields prediction intervals
- Output: $(\hat{Y}_t, \hat{\sigma}_t)$ — point estimate with uncertainty

Results: Estimation Task (Spearman ρ)

Model	DE		FR		DE_FR	
	d=2	d=3	d=2	d=3	d=2	d=3
DS3M_UV	.78 $\pm$.38	.58 $\pm$.43	.69 $\pm$.40	.93 $\pm$.02	.45 $\pm$.32	.25 $\pm$.34
DS3M_MV	.41 $\pm$.46	.58 $\pm$.49	.75 $\pm$.37	.44 $\pm$.42	.34 $\pm$.40	.24 $\pm$.27
FANTOM_BEM	.61 $\pm$.20	.92 $\pm$.05	.83 $\pm$.18	.89 $\pm$.08	.55 $\pm$.39	.69 $\pm$.14
CaRS	.74 $\pm$.06	.72 $\pm$.06	.43 $\pm$.12	.47 $\pm$.08	.44 $\pm$.11	.47 $\pm$.19

Key Findings:

- **DE:** FANTOM_BEM achieves $\rho = 0.92$ (d=3), highest among all models
- **FR:** DS3M_UV strong ($\rho = 0.93$) for d=3; FANTOM_BEM best for d=2 (.83)
- **DE_FR:** FANTOM_BEM best ($\rho = 0.55$ – 0.69)
- **CaRS:** Most stable across seeds ($\sigma = 0.06$ – 0.19 vs DS3M $\sigma = 0.27$ – 0.49)

Results: Estimation Task (RMSE)

Model	DE		FR		DE_FR	
	d=2	d=3	d=2	d=3	d=2	d=3
DS3M_UV	190 $\pm_{10}$	199 $\pm_{11}$	103 $\pm_{15}$	96 $\pm_{10}$	1375 $\pm_{20}$	1412 $\pm_{69}$
DS3M_MV	199 $\pm_{14}$	201 $\pm_6$	103 $\pm_{12}$	108 $\pm_{14}$	1395 $\pm_{32}$	1386 $\pm_{31}$
FANTOM_BEM	12883 $\pm_{10852}$	21580 $\pm_{11157}$	4508 $\pm_{1971}$	4110 $\pm_{1953}$	46187 $\pm_{21589}$	81941 $\pm_{111418}$
CaRS	205 $\pm_1$	205 $\pm_1$	117 $\pm_2$	116 $\pm_1$	1370 $\pm_3$	1376 $\pm_6$

RMSE in EUR/MWh for DE/FR, spread units for DE_FR.

Key Findings:

- **DS3M models:** Comparable RMSE (~ 190 - 206 for DE, ~ 100 - 117 for FR)
- **FANTOM_BEM:** Very high RMSE due to different prediction scale (raw price vs returns)
- **CaRS:** Most stable RMSE with very low variance (± 1 - 14)

Results: Estimation Task (Directional Accuracy)

Model	DE		FR		DE_FR	
	d=2	d=3	d=2	d=3	d=2	d=3
DS3M_UV	.84 $\pm$.19	.72 $\pm$.23	.75 $\pm$.23	.90 $\pm$.02	.65 $\pm$.11	.64 $\pm$.11
DS3M_MV	.68 $\pm$.16	.75 $\pm$.24	.80 $\pm$.15	.68 $\pm$.21	.56 $\pm$.07	.59 $\pm$.11
FANTOM_BEM	.51 $\pm$.04	.52 $\pm$.03	.51 $\pm$.04	.49 $\pm$.04	.50 $\pm$.01	.52 $\pm$.07
CaRS	.73 $\pm$.05	.66 $\pm$.11	.63 $\pm$.06	.61 $\pm$.08	.64 $\pm$.05	.67 $\pm$.08

Directional accuracy = fraction of correct price movement direction predictions.

Key Findings:

- **FR d=3:** DS3M_UV achieves 90% directional accuracy – strong autoregressive signal
- **FANTOM_BEM:** Near-random (50%) directional accuracy despite reasonable Spearman
- **CaRS:** Consistent 61–73% accuracy ($\sigma = 0.05$ – 0.11 vs DS3M $\sigma = 0.15$ – 0.24)

Results: Prediction Task (Spearman ρ)

Model	DE		FR		DE_FR	
	d=2	d=3	d=2	d=3	d=2	d=3
DS3M_UV	.27 $\pm$.12	.13 $\pm$.22	.25 $\pm$.17	.08 $\pm$.15	.11 $\pm$.07	.11 $\pm$.10
DS3M_MV	.29 $\pm$.26	.27 $\pm$.23	.33 $\pm$.25	.25 $\pm$.24	.11 $\pm$.11	.07 $\pm$.08
FANTOM_BEM	.82 $\pm$.18	.93 $\pm$.03	.81 $\pm$.18	.93 $\pm$.02	.64 $\pm$.09	.69 $\pm$.11
CaRS	.45 $\pm$.05	.44 $\pm$.09	.42 $\pm$.07	.44 $\pm$.02	.17 $\pm$.05	.17 $\pm$.05

Key Findings:

- **FANTOM_BEM dominates:** Very high ρ (0.81–0.93) across all configurations
- **But:** High Spearman with $\sim 50\%$ directional accuracy suggests overfitting to trends
- **CaRS:** Moderate Spearman (0.17–0.45) but better directional accuracy (55–65%)

Results: Prediction Task (RMSE)

Model	DE		FR		DE_FR	
	d=2	d=3	d=2	d=3	d=2	d=3
DS3M_UV	196 ± 7	205 ± 5	112 ± 5	115 ± 6	1419 ± 24	1419 ± 11
DS3M_MV	202 ± 7	204 ± 3	114 ± 4	114 ± 5	1416 ± 11	1419 ± 25
FANTOM_BEM	6837 ± 7818	25130 ± 14301	4630 ± 1439	2963 ± 1077	72240 ± 85922	51741 ± 26566
CaRS	205 ± 1	206 ± 1	116 ± 1	116 ± 0	1392 ± 3	1394 ± 4

RMSE in EUR/MWh for DE/FR, spread units for DE_FR.

Key Findings:

- **DS3M models:** Similar RMSE (~ 196 - 206 for DE, ~ 112 - 116 for FR)
- **CaRS:** Most stable RMSE with very low variance (± 1 - 4)
- **FANTOM_BEM:** High RMSE variance indicates unstable predictions

Results: Prediction Task (Directional Accuracy)

Model	DE		FR		DE_FR	
	d=2	d=3	d=2	d=3	d=2	d=3
DS3M_UV	.57 $\pm$.06	.50 $\pm$.08	.56 $\pm$.08	.52 $\pm$.06	.52 $\pm$.02	.52 $\pm$.01
DS3M_MV	.60 $\pm$.09	.58 $\pm$.10	.58 $\pm$.10	.57 $\pm$.09	.51 $\pm$.01	.51 $\pm$.01
FANTOM_BEM	.50 $\pm$.08	.54 $\pm$.00	.49 $\pm$.04	.52 $\pm$.04	.50 $\pm$.01	.50 $\pm$.00
CaRS	.65$\pm$.02	.62$\pm$.02	.64$\pm$.03	.64$\pm$.02	.55$\pm$.02	.55$\pm$.02

Key Findings:

- **CaRS**: Best directional accuracy for DE/FR (62-65%); best for DE_FR (.55)
- **DS3M_MV**: Competitive performance on DE (58-60%)
- **FANTOM_BEM**: Near-random ($\sim$ 50%) despite high Spearman – optimizes for correlation, not direction

Estimation vs Prediction Gap (Directional Accuracy)

	DE		FR		DE_FR	
	Est	Pred	Est	Pred	Est	Pred
CaRS (d=2)	.73 $\pm$.05	.65 $\pm$.02	.63 $\pm$.06	.64 $\pm$.03	.64 $\pm$.05	.55 $\pm$.02
Δ		-0.08		+0.01		-0.09

Directional accuracy = fraction of correct price movement direction predictions.

Insights:

- **FR:** Prediction *matches* estimation ($\Delta \approx 0$) – autoregressive signal generalizes well
- **DE:** Small gap ($\Delta = -0.08$) – CaRS maintains 65% accuracy for prediction
- **DE_FR:** Largest gap ($\Delta = -0.09$) – cross-country dynamics harder to predict
- **Key insight:** CaRS shows **stable directional accuracy** across tasks (unlike Spearman which degrades more)

Discussion: Why Different Models Excel at Different Tasks

DS3M Dominates Estimation

- Electricity prices **highly autocorrelated**: DE $\rho_{\text{lag-1}} = 0.94$, FR $\rho_{\text{lag-1}} = 0.97$
- DS3M's GRU captures this **autoregressive signal**
- Result: Up to 90% directional accuracy (FR $d=3$)
- **But**: High variance ($\sigma = 0.19\text{--}0.49$) vs CaRS ($\sigma = 0.02\text{--}0.12$)

Why autocorrelation helps estimation:

- Yesterday's price strongly predicts today's
- DS3M_UV uses *only* target history—perfect for exploiting temporal persistence

CaRS Excels at Prediction

- One-step-ahead: autoregressive signal **degrades**
- DS3M_UV Dir.Acc drops to 50–57% (near random)
- CaRS discovers **causal relationships** that:
 - Generalize beyond temporal patterns
 - Adapt via regime-specific DAGs

Why causal structure helps prediction:

- Causal relationships are more **stable** across time
- Not dependent on short-term autocorrelation
- Regime switching captures structural breaks

Key Insight: Strong autocorrelation ($\rho > 0.9$) is a *double-edged sword*—excellent for in-sample estimation, but the signal degrades out-of-sample where causal structure provides more robust predictions.

DS3M Baseline Variance

- DS3M: $\sigma = 0.19\text{--}0.49$ (Spearman) vs CaRS: $\sigma = 0.02\text{--}0.12$ (5–10 $\times$ lower)
- Some configurations: $\sigma > 50\%$ of the mean
- CaRS's regularization (DAG + sparsity) acts as implicit **training stabilizer**

FANTOM_BEM Scale Mismatch

- Predicts **returns**, not price levels
- Inflates Spearman ρ but Dir.Acc $\approx 50\%$ (random)
- RMSE 12k–82k vs DS3M/CaRS ~ 200
- **Spearman not directly comparable**

Estimation vs Prediction Tradeoff

- DS3M's estimation advantage is largely an **autocorrelation artifact** ($\rho_{\text{lag-1}} = 0.94\text{--}0.97$)
- CaRS trades some estimation for causal interpretability, prediction robustness, and training stability
- For **trading**: prediction stability (low σ , consistent Dir.Acc) matters more than peak estimation Spearman

Takeaway: DS3M excels when autocorrelation is exploitable. CaRS excels when **interpretability** and **robustness** are required.

Motivation:

- Standard CaRS learns **independent DAGs** per regime
- **Shared Backbone** learns:
 - Common causal structure across regimes
 - Regime-specific edge modifications
- Hypothesis: Shared structure improves generalization

Implementation:

- Base DAG A_{shared} learned jointly
- Per-regime adjustments $\Delta A^{(d)}$
- Final DAG: $A^{(d)} = A_{\text{shared}} + \Delta A^{(d)}$

Experimental Setup:

- Markets: DE, FR, DE_FR
- Regimes: $d=2$, $d=3$
- Seeds: 42, 123, 456, 789, 1011
- Total: 30 experiments
- Task: Prediction

Key Differences from CaRS:

- Enforces structural similarity
- Reduces parameter count
- May improve stability

Shared Backbone: Results Summary

Metric	DE		FR		DE_FR	
	d=2	d=3	d=2	d=3	d=2	d=3
Spearman ρ	$.37 \pm .10$	$.31 \pm .12$	$.28 \pm .10$	$.29 \pm .18$	$.02 \pm .03$	$.07 \pm .02$
RMSE	$1.41 \pm .01$	$1.41 \pm .01$	$2.59 \pm .04$	$2.56 \pm .06$	$1.41 \pm .00$	$1.42 \pm .01$
Dir. Accuracy	$.75 \pm .06$	$.75 \pm .03$	$.75 \pm .01$	$.72 \pm .04$	$.75 \pm .01$	$.70 \pm .05$

Results aggregated over 5 seeds. Dir. Accuracy = fraction of correct price direction predictions.

Key Findings:

- **Directional accuracy significantly higher** than standard CaRS (70-75% vs 55-65%)
- **DE_FR spread prediction:** 75% directional accuracy despite near-zero Spearman ($\rho \approx 0.02$)
- Shared structure appears to improve directional prediction capability

Shared Backbone: Estimation Task Results

Metric	DE		FR		DE_FR	
	d=2	d=3	d=2	d=3	d=2	d=3
Spearman ρ	.51 $\pm$.17	.41 $\pm$.06	.36 $\pm$.09	.35 $\pm$.18	.02 $\pm$.01	.03 $\pm$.04
RMSE	1.40 $\pm$.00	1.40 $\pm$.01	2.58 $\pm$.02	2.55 $\pm$.06	1.40 $\pm$.00	1.42 $\pm$.00
Dir. Accuracy	.78 $\pm$.06	.74 $\pm$.07	.76 $\pm$.01	.74 $\pm$.03	.76 $\pm$.00	.66 $\pm$.02

Estimation task: Results aggregated over 5 seeds (42, 123, 456, 789, 1011).

Key Findings:

- **Directional accuracy 74-78%** for DE and FR (vs 70-75% for prediction)
- **DE:** Highest Spearman ($\rho = 0.51$ for d=2) with 78% directional accuracy
- **DE_FR d=2:** Strong directional accuracy (76%) despite near-zero Spearman
- **Consistent performance** across seeds (low variance)

Shared Backbone: Estimation vs Prediction Comparison

Task (d=2)	DE		FR		DE_FR	
	Est	Pred	Est	Pred	Est	Pred
Spearman ρ	.51 $\pm$.17	.37 $\pm$.10	.36 $\pm$.09	.28 $\pm$.10	.02 $\pm$.01	.02 $\pm$.03
Dir. Accuracy	.78 $\pm$.06	.75 $\pm$.06	.76 $\pm$.01	.75 $\pm$.01	.76 $\pm$.00	.75 $\pm$.01
Δ Dir. Acc.	+0.03		+0.01		+0.01	

Key Result: Shared Backbone maintains high directional accuracy (74-78%) for both estimation and prediction tasks, with only marginal improvement for estimation.

Interpretation:

- Estimation benefits from access to concurrent features X_t
- Small gap suggests model generalizes well to prediction task
- Shared causal structure provides stability across both tasks

Shared Backbone vs CaRS: Directional Accuracy Comparison (Prediction)

Model	DE		FR		DE_FR	
	d=2	d=3	d=2	d=3	d=2	d=3
CaRS	.65 $\pm$.02	.62 $\pm$.02	.64 $\pm$.03	.64 $\pm$.02	.55 $\pm$.02	.55 $\pm$.02
Shared Backbone	.75$\pm$.06	.75$\pm$.03	.75$\pm$.01	.72$\pm$.04	.75$\pm$.01	.70$\pm$.05
Δ	+0.10	+0.13	+0.11	+0.08	+0.20	+0.15

Key Result: Shared Backbone achieves **+8 to +20 percentage points** higher directional accuracy than standard CaRS across all markets and regime configurations.

Interpretation:

- Shared causal structure acts as **regularization**
- Prevents overfitting to regime-specific noise
- Especially beneficial for cross-market prediction (DE_FR: +20pp)

Top Performers by Market:

DE (Germany):

- Best: d3_seed123 (79.1% dir. acc.)
- Runner-up: d2_seed42 (78.9%)
- Spearman range: 0.16–0.48

FR (France):

- Best: d3_seed789 (77.5% dir. acc.)
- Runner-up: d2_seed456 (76.5%)
- Spearman range: -0.02–0.44

DE_FR (Spread):

- Best: d2_seed123 (76.1% dir. acc.)
- Runner-up: d2_seed42 (76.0%)
- Spearman range: -0.01–0.09

Trading Implications:

Directional Trading Strategy:

With 75% directional accuracy:

- Win rate: 3 out of 4 trades
- Edge over random: +25pp
- Suitable for trend-following

Note on Spearman vs Direction:

- Low Spearman ($\rho \approx 0$) can coexist with high directional accuracy
- Model predicts *direction* better than *magnitude*
- For trading, direction often matters more

Results: Edge Statistics and Regime Detection

Learned DAG Sparsity — FANTOM:

Dataset	d	$\lambda=1$	$\lambda=5$	$\lambda=10$	$\lambda=50$	Max*
DE	2	2465	1410	596	247	11025
DE	3	3142	1364	772	253	11025
FR	2	2255	1427	1092	275	7225
FR	3	2495	1360	491	219	7225
DE_FR	2	391	298	243	129	1024
DE_FR	3	396	299	228	120	1024

*Max = n^2 (lag=1). Avg non-zero edges per regime, mean across seeds.

DS3MCausal produces near-saturated graphs (>95% of max) for all λ .

Sparsity: 1–22% (FANTOM), ~2% at $\lambda_{\text{sparse}}=50$

Regime Collapse Issue:

- >95% of d=4 runs collapse to 1-2 regimes
- d=2,3 also show frequent collapse
- **Solution:** Report d=2,3 only

Successful 2-Regime Detection:

- DE (seed 42): 722/2929 split
- FR (seed 42): 697/2954 split
- DE-FR (seed 42): 3302/238 split
- Regimes capture different market conditions

Issue: Regime collapse reduces interpretability benefit.
d=4 excluded from results due to >95% collapse rate.

Edge Count by Lag Setting:

Dataset	lag=1	lag=2	lag=3
DE	737	1294	1756
FR	681	1335	1766
DE_FR	1017	1644	2362

Values = mean total edges (all weights) across regimes and seeds

Observation: Edges scale $\sim$ linearly with lag (more temporal connections discovered).

Sparsity vs. λ_{sparse} :

- $\lambda = 1.0$: Dense graphs (many weak edges)
- $\lambda = 5.0$: Moderate sparsity (balanced)
- $\lambda = 10.0+$: Sparse graphs (few strong edges)

Top Causal Features by Regime (DE,

$d=2$):

	Regime 0	Regime 1
1	Net Flow to FR	Actual Load
2	Fossil Gas	Fossil Gas
3	Wind Onshore	Day-Ahead Price
4	Wind Speed (10m)	Wind Speed (100m)
5	Natural Gas	Hydro Pumped

Top Causal Features by Regime (FR,

$d=2$):

	Regime 0	Regime 1
1	Wind Offshore	Wind Speed (10m)
2	Hydro Pumped	Solar Radiation
3	Precipitation	Hydro Pumped
4	Cloud Cover	Nuclear
5	Nuclear	Wind Offshore

Insight: CaRS discovers economically plausible causal drivers (generation sources $\rightarrow$ price).

Regime 0

Regime 1

Legend: Blue = instantaneous edges (t), Orange dashed = lagged edges ($t-1 \rightarrow t$)

Regime 0

Regime 1

Legend: Blue = instantaneous edges (t), Orange dashed = lagged edges ($t-1 \rightarrow t$)

Are Regimes Comparable Across K Settings?

Challenge: Regime Interpretation

- Regimes learned for $K = 2$ are **not directly comparable** to those for $K = 3$ or $K = 4$
- Each K setting learns a different partitioning of the data
- No ground truth regime labels available

Observed Patterns:

- $K = 2$: Often separates normal vs. crisis periods (2022 energy crisis)
- $K = 3, 4$: May capture finer granularity (seasonal, intraday)
- Higher $K \rightarrow$ more regime collapse risk

Regime Distribution Examples:


Setting	Seed	Regime Split
DE, $K = 2$	42	722 / 2929
FR, $K = 2$	42	697 / 2954
FR, $K = 3$	123	2104 / 4821 / 3099

Recommendation:

- Start with $K = 2$ for interpretability
- Increase K only if evidence suggests multiple distinct dynamics
- Validate regimes against known market events


What Drives Price Changes? Regime-Specific DAGs

Regime 0 (Normal Market):



Drivers: Renewables + Load

Regime 1 (Crisis/High Volatility):



Drivers: Fossil fuels + Carbon

Legend: ● Target (Price) ● Renewable/Load ● Commodity (crisis driver)
→ Strong causal → Crisis-dominant - -> Weak/suppressed

What Drives Price Changes? Regime-Specific DAGs

Interpretation:

- **Normal regime:** Renewables (wind, solar) and load dominate price formation
- **Crisis regime:** Commodity prices (gas, carbon) become primary drivers

2021-2022 Energy Crisis (Documented)

- Volatility increased **159% in DE, 185% in FR** (2021 vs 2020)
- Prices: 35 EUR/MWh (2020) → 500+ EUR/MWh (Mar 2022)
- “A clear dividing line between stable pre-crisis and new era of increased volatility” [Keles & Dehler-Holland, 2025]

France-Specific Dynamics

- Strongest price increase despite no Russian gas dependence
- Nuclear unavailability drove regime change [Witthaut et al., 2024]

MRS Models for Electricity (Established)

- Huisman & Mahieu (2003): First MRS application to DE/UK markets
- “MRS models are a natural candidate for modeling spiky behavior” [Janczura & Weron, 2010]
- Recent: Kapoor & Balakrishna (2023) use 2-5 regimes for forecasting

Our Results Align With:

- ✓ Pre-2021 = Normal regime (low volatility)
- ✓ 2021+ = Crisis regime (high volatility)
- ✓ Gas/commodities as crisis drivers

How Does the Setting Matter?

Task: Estimation vs Prediction

- **Estimation** ($Y_t | X_{0:t}, Y_{0:t-1}$):
 - Higher accuracy (access to recent Y)
 - Primarily autoregressive signal
 - Univariate models excel
- **Prediction** ($Y_{t+1} | X_{0:t}, Y_{0:t}$):
 - More challenging
 - CaRS shows **best directional accuracy** (62-65%)
 - Causal structure helps generalize

Target: DE vs FR vs DE_FR

- **FR**: Highly autoregressive
 - DS3M_UV achieves $\rho = 0.93$ ($d=3$)
 - Nuclear baseload = stable dynamics
- **DE**: More complex
 - Higher renewable share
 - CaRS best for directional prediction
- **DE_FR**: Joint modeling hardest
 - CaRS achieves 55% directional accuracy
 - Best among all models

Recommendations:

DS3M_UV: Best Spearman for FR (autoregressive)

CaRS: Best interpretability for causal analysis

Shared Backbone: Best directional accuracy (70-75%) for trading applications

Avoid FANTOM_BEM: Scalability issues (>50 features $\rightarrow$ OOM on 80GB GPU)

1. Seed Initialization Variance

- High variance across seeds ($\sigma > 0.1$)
- **Solutions:**
 - Better initialization schemes
 - Ensemble methods
 - Bayesian treatment of DAG parameters

2. Regime Collapse

- Often converges to 1 regime
- Need stronger regime regularization
- Entropy bonus for regime diversity

3. Additional Data Sources

- **TTF Gas:** European natural gas benchmark
- **API2 Coal:** Coal price index
- **EU ETS Carbon:** Carbon allowance prices
- Weather forecasts, demand forecasts

4. Model Enhancements

- **Ensembles:** Combine multiple seeds for robust predictions
- **Longer lags:** Currently lag=1, may benefit from lag=2+

Key Challenge: Balancing model complexity (more regimes, more edges) with stability and interpretability.

- **CaRS** successfully combines regime switching (DS3M) with causal discovery (FANTOM)
- **Shared Backbone variant** achieves best directional accuracy:
 - Prediction: DE **75%**, FR **72-75%**, DE_FR **70-75%**
 - Estimation: DE **78%**, FR **76%**, DE_FR **76%** (d=2)
 - Consistent across both tasks (small estimation-prediction gap)
- **Key insight:** Shared causal structure improves directional prediction by +8 to +20pp
- **Trade-offs:**
 - DS3M_UV: Best Spearman for autoregressive tasks (FR)
 - CaRS: Best interpretability with moderate directional accuracy
 - **Shared Backbone: Best directional accuracy for trading applications**
- **Open challenges:** Regime collapse, seed variance, Spearman-direction disconnect

Questions?

Regime-Switching Models:

- Hamilton (1989): Markov-switching autoregressive models for business cycles
- Kim (1994): Dynamic linear models with regime switching
- Xu et al. (2021): **DS3M** – Deep Switching State Space Model

Causal Discovery Methods:

- Spirtes et al. (2000): PC algorithm for constraint-based discovery
- Chickering (2002): GES – Greedy Equivalence Search
- Zheng et al. (2018): **NOTEARS** – Differentiable DAG learning
- Ng et al. (2020): **GOLEM** – Likelihood-based structure learning
- Pamfil et al. (2020): **DYNOTEARS** – Time series extension

Deep Causal Models:

- Geffner et al. (2022): **DECI** – Deep End-to-end Causal Inference
- Gong et al. (2022): **Rhino** – Historically dependent noise
- Rahmani et al. (2025): **FANTOM** – Regime-switching causal discovery

Energy Price Modeling:

- Weron (2014): Electricity price forecasting – comprehensive review
- Ziel & Weron (2018): Day-ahead electricity price forecasting
- Lago et al. (2021): Deep learning for electricity price forecasting
- Marcjasz et al. (2023): Probabilistic electricity price forecasting

State Space Models:

- Kalman (1960): Linear state estimation
- Durbin & Koopman (2012): Time series analysis by state space methods
- Krishnan et al. (2017): Structured inference networks for nonlinear state space
- Rangapuram et al. (2018): Deep state space models for time series forecasting

Variational Inference:

- Kingma & Welling (2014): VAE – Variational Auto-Encoders
- Rezende et al. (2014): Stochastic backpropagation
- Jang et al. (2017): Gumbel-Softmax for discrete latent variables

Energy Crisis & Regime Validation:

- Keles & Dehler-Holland (2025): Volatility analysis 2015-2025, MDPI Applied Sciences
- Gonzalez-Aparicio et al. (2024): North Sea electricity markets 2018-2022, Energy Economics
- Witthaut et al. (2024): Network analysis of 2021-22 crisis, Chaos Journal
- Gil-Tertre et al. (2023): Structural changes in energy markets, ECB Forum

Markov Regime-Switching for Electricity:

- Huisman & Mahieu (2003): First MRS for electricity prices (DE/UK)
- Mount et al. (2006): Predicting price spikes with time-varying MRS, Energy Economics
- Janczura & Weron (2010): Efficient MRS estimation, Computational Statistics
- Kapoor & Balakrishna (2023): MRS forecasting (2-5 regimes), Journal of Forecasting

Key Finding: Our detected Normal/Crisis regimes align with documented structural break in European electricity markets starting 2021.

Appendix: Regime Time Series Visualization (DE)

Appendix: Regime Time Series Visualization (FR)

Appendix: Regime Time Series Visualization (DE-FR)

Appendix: Causal Feature Importance (DE)

Appendix: Causal Feature Importance (FR)

Appendix: Causal Feature Importance (DE-FR)

Appendix: Lag Analysis

Edge count increases linearly with lag (more temporal connections discovered).

Appendix: Edge Statistics Summary

Metric	Germany (DE)		France (FR)	
	Regime 0	Regime 1	Regime 0	Regime 1
Sample Count	722 (20%)	2929 (80%)	697 (19%)	2954 (81%)
<i>Instantaneous Edges (lag=0):</i>				
# Edges	420	420	272	272
Avg Weight	0.000113	0.000115	0.000124	0.000126
Max Weight	0.000312	0.000462	0.000254	0.000294
<i>Lagged Edges (lag=1):</i>				
# Edges	441	441	289	289
Avg Weight	1.21e-06	1.25e-06	2.22e-06	2.22e-06
Max Weight	3.14e-06	5.54e-06	4.60e-06	4.73e-06

Table : Edge statistics from DS3MCausal with native Markov regime switching ($d=2$, $\text{lag}=1$, $\lambda_s=5.0$, $\text{seed}=42$)

Key observation: Germany Regime 1 shows significantly higher max instantaneous weight (0.000462), indicating more concentrated causal relationships during crisis periods.

Appendix: Germany Feature Importance by Regime

Regime 0 (Normal Market)

Feature	Total Weight
Net Flow FR	0.00498
Fossil Gas	0.00494
Wind Onshore	0.00493
Wind 10m	0.00485
Nat Gas Price	0.00481
Temperature	0.00464
Lignite	0.00463

Renewables + cross-border flows dominant

Regime 1 (Crisis/High Volatility)

Feature	Total Weight
Actual Load	0.00577
Fossil Gas	0.00524
Day Ahead Price	0.00515
Wind 100m	0.00473
Hydro Pumped	0.00471
Net Flow FR	0.00462
Wind Onshore	0.00460

Load becomes dominant driver (+47% vs Regime 0)

Source: cars_de_d2_lag1_ls5.0_seed42_ds3m_causal_native_markov.json

Appendix: Germany Instantaneous Edges TO Price

Feature	Regime 0	Regime 1	Change
Actual Load	0.000107	0.000462	+332%
Lignite	0.0000986	0.000174	+77%
Hydro Pumped	0.000140	0.000119	-15%
Solar	0.000130	0.000116	-11%
Nat Gas Price	0.000129	0.000102	-21%
Net Flow FR	0.000123	0.000107	-13%
Fossil Gas	0.000104	0.000101	-3%
Wind Onshore	0.000104	0.0000944	-9%
Wind Offshore	0.000103	0.000109	+6%
Nuclear	0.0000989	0.0000994	+0.5%

Table : Instantaneous causal edges TO Price (adjacency matrix column 0)

Interpretation: During crisis, electricity **demand (Load)** becomes the overwhelming price driver ($4.3\times$ higher than any other edge), while renewables' influence diminishes.

Appendix: France Feature Importance by Regime

Regime 0 (Normal Market)

Feature	Total Weight
Wind Offshore	0.00431
Hydro Pumped	0.00422
FR Precipitation	0.00421
Cloud Cover	0.00419
Nuclear	0.00416
Wind Onshore	0.00408
Temperature	0.00405

Weather + nuclear dominant

Regime 1 (Crisis/High Volatility)

Feature	Total Weight
Wind 10m	0.00453
FR Radiation	0.00434
Hydro Pumped	0.00424
Nuclear	0.00422
Wind Offshore	0.00419
Fossil Gas	0.00415
Actual Load	0.00412

Weather features increase in importance

Source: cars_fr_d2_lag1_ls5.0_seed42_ds3m_causal_native_markov.json

Appendix: France Instantaneous Edges TO Price

Feature	Regime 0	Regime 1	Change
Solar	0.000117	0.000161	+38%
Nat Gas Price	0.000116	0.000132	+14%
Cloud Cover	0.000107	0.000123	+15%
Actual Load	0.000156	0.000121	-22%
Wind Offshore	0.000139	0.000110	-21%
Wind 100m	0.000119	0.0000983	-17%
Hard Coal	0.000131	0.000119	-9%
Hydro Pumped	0.000127	0.000123	-3%
Nuclear	0.000117	0.000111	-5%
Fossil Gas	0.000125	0.000121	-3%

Table : Instantaneous causal edges TO Price (adjacency matrix column 0)

Interpretation: France shows a more balanced regime shift – Solar becomes dominant in crisis regime (+38%), while Load decreases in importance (-22%). More subtle structural change than Germany.

Appendix: Cross-Country Comparison

Metric	Germany (DE)	France (FR)
<i>Regime Distribution:</i>		
Regime 0 (Normal)	722 samples (20%)	697 samples (19%)
Regime 1 (Crisis)	2929 samples (80%)	2954 samples (81%)
<i>Regime 0 Top Driver:</i>		
Feature	Hydro Pumped	Actual Load
Weight	0.000140	0.000156
<i>Regime 1 Top Driver:</i>		
Feature	Actual Load	Solar
Weight	0.000462	0.000161
<i>Regime Differentiation:</i>		
Max weight ratio (R1/R0)	4.3×	1.0×
Top driver change?	Yes (Hydro → Load)	Yes (Load → Solar)

Key Finding: Germany shows **dramatic regime differentiation** (Load ↑ 332% in crisis), while France shows **more balanced regimes**. This aligns with Germany's higher gas dependency during the 2021-2022 energy crisis.

Appendix: FANTOM-BEM with Price Returns Target

Motivation: FANTOM-BEM originally predicts raw Day Ahead Price (EUR/MWh), while CaRS/DS3M predict price_change_pct (returns). We re-ran FANTOM-BEM with returns as the target variable to ensure comparable evaluation.

Market	d=2 Collapse	d=3 Collapse	Edge Count	Top Feature
DE	1/5 (20%)	5/5 (100%)	127–149	price_change_pct
FR	5/5 (100%)	5/5 (100%)	82–85	price_change_pct
DE_FR	5/5 (100%)	5/5 (100%)	122–124	price_spread_change_pct
Overall	83% Regime Collapse		–	–

Table : Regime collapse rates when using price returns as target (5 seeds per config)

Key Finding: Using returns instead of raw prices causes **massive regime collapse** (83% overall). This reveals that the original regime detection was driven by **price level autocorrelation**, not genuine market regime dynamics.

Implication: Raw prices exhibit strong autocorrelation that creates spurious temporal structure. Returns remove this

Raw Prices (Original)

- Strong autocorrelation ($\rho > 0.9$)
- Price levels persist across days
- Regime detection captures price “stickiness”
- Distinct regime structure (20/80% split)

Price Returns (New)

- Near-zero autocorrelation
- Day-to-day changes are more random
- Regime detection loses signal
- Regimes collapse to single state

Metric	Raw Prices	Returns
Regime Collapse Rate	<5%	83%
Regime Differentiation	Strong ($4.3\times$)	None
RMSE Scale	12,000–80,000	~100–300
Comparable to CaRS?	No	Yes (scale only)

Conclusion: The regime detection in FANTOM-BEM relies on price level dynamics. When switching to returns, the model cannot identify meaningful regimes, suggesting the Bayesian EM component captures *price persistence patterns* rather than *market state transitions*.

Appendix: CaRS DAG Learning Ablation Study

Challenge: The shared backbone architecture produces identical DAGs across regimes.

Ablation experiments (5 seeds each, mean $\pm$ std):

Experiment	Directional Accuracy (%)			DAG L2 Diff
	DE	FR	DE-FR	
Independent	74.9 $\pm$ 3.9	74.3 $\pm$ 1.5	75.1 $\pm$ 1.5	0.0
High λ_{diff}	77.0$\pm$3.9	74.9$\pm$1.3	75.2$\pm$1.0	0.0
Noise Init	74.3 $\pm$ 4.3	74.0 $\pm$ 3.7	72.9 $\pm$ 4.5	3.6
Combined	74.4 $\pm$ 4.3	74.0 $\pm$ 3.7	72.0 $\pm$ 4.2	3.6

Experiment	Regime Balance (R0/R1)			Spearman Correlation		
	DE	FR	DE-FR	DE	FR	DE-FR
Independent	266/449	406/309	429/286	0.30	0.19	0.01
High λ_{diff}	272/443	412/303	429/286	0.41	0.28	0.01
Noise Init	313/402	335/381	153/562	0.32	0.30	0.00
Combined	313/402	335/381	147/569	0.32	0.30	0.02

Key Findings:

- **High λ_{diff} :** Best accuracy on *all* markets (+2.1% on DE), most stable
- **Noise init:** Creates DAG differentiation (L2=3.6) but causes regime collapse on DE-FR
- **Trade-off:** DAG differentiation $\leftrightarrow$ accuracy and regime stability

Noise initialization creates regime-specific DAG structures (seed 42 example):

DE Market – Noise Init

Metric	Regime 0	Regime 1
Mean edge weight	0.468	0.480
Max edge weight	0.888	0.998
Edges > 0.9	0	48
Edges TO Price (max)	0.864	0.942

FR Market – Noise Init

Metric	Regime 0	Regime 1
Mean edge weight	0.601	0.615
Max edge weight	0.901	0.988
Edges > 0.9	1	49
Edges TO Price (max)	0.864	0.870

Interpretation:

- Regime 1 develops stronger, more binary edges (closer to 0/1)
- DAG L2 difference: 3.6 ± 0.2 (DE), 3.7 ± 0.5 (FR) across 5 seeds
- **Trade-off**: DAG differentiation comes at cost of accuracy (-2.7% on DE)
- **Warning**: Causes regime collapse on DE-FR (153/562 vs 429/286 for high λ)

Why does shared backbone fail to differentiate regimes?

Shared Backbone Design:

$$A_{\text{regime}_d} = A_{\text{shared}} \oplus A_{\text{deviation}_d}$$

where $\oplus$ is probabilistic OR.

Problem:

- Deviations initialized to zero
- Gradients from L2 penalty too weak
- Both regimes converge to shared backbone
- Cosine similarity stays at 1.0

Solutions Tested:

- 1 ✓ Higher λ_{diff} (10 $\times$)
- 2 ✓ Noise initialization
- 3 ✓ Independent DAGs


Recommendation:

Use **high** λ_{diff} for best directional accuracy across all markets (77.0% on DE).

Key Insight: High λ_{diff} achieves best accuracy (77% on DE) while maintaining regime stability. Noise initialization creates DAG differentiation but degrades performance.


Native Markov / Independent Mode

- **NOT hierarchical** – flat structure
- K completely independent DAGs
- No shared structure across regimes
- Each regime: A_d learned independently



Shared Backbone Mode

- **Hierarchical** two-level structure
- Global meta DAG: A_{shared}
- Regime-specific residuals: ΔA_d
- Combined:
$$A_d = 1 - (1 - A_{\text{shared}})(1 - \Delta A_d)$$



ICGNN: Both configurations use the **same CausalICGNN class** with K instances (one per regime).

Appendix: ELBO Derivation

Starting Point: We want to maximize the log-likelihood $\log p(Y|X)$, but it's intractable.

Step 1: Introduce variational distribution q

$$\log p(Y|X) = \log \int p(Y, z, d, G|X) dz = \log \int q(z, d, G) \frac{p(Y, z, d, G|X)}{q(z, d, G)} dz$$

Step 2: Apply Jensen's Inequality (since log is concave)

$$\log p(Y|X) \geq \int q(z, d, G) \log \frac{p(Y, z, d, G|X)}{q(z, d, G)} dz = \mathbb{E}_q \left[\log \frac{p(Y, z, d, G|X)}{q(z, d, G)} \right]$$

Step 3: Expand the ELBO

$$\text{ELBO} = \mathbb{E}_q[\log p(Y|X, z, d, G)] + \mathbb{E}_q[\log p(z, d, G)] - \mathbb{E}_q[\log q(z, d, G)]$$

Step 4: Recognize the gap

$$\log p(Y|X) = \text{ELBO} + \underbrace{\text{KL}[q(z, d, G) \parallel p(z, d, G|Y, X)]}_{\geq 0}$$

Gaussian KL (closed form):

For $p = \mathcal{N}(\mu_1, \sigma_1^2)$ and $q = \mathcal{N}(\mu_2, \sigma_2^2)$:

$$\text{KL}[p\|q] = \log \frac{\sigma_2}{\sigma_1} + \frac{\sigma_1^2 + (\mu_1 - \mu_2)^2}{2\sigma_2^2} - \frac{1}{2}$$

For standard normal prior $q = \mathcal{N}(0, 1)$:

$$\text{KL}[\mathcal{N}(\mu, \sigma^2)\|\mathcal{N}(0, 1)] = -\frac{1}{2}(1 + \log \sigma^2 - \mu^2 - \sigma^2)$$

Categorical KL:

For $p = \text{Cat}(\pi)$ and $q = \text{Cat}(\rho)$:

$$\text{KL}[p\|q] = \sum_{k=1}^K \pi_k \log \frac{\pi_k}{\rho_k}$$

Application to CaRS:

- $\text{KL}(z_t)$: Uses Gaussian formula with $\mu(h_t, d_t)$, $\sigma(h_t, d_t)$
- $\text{KL}(d_t)$: Categorical KL between posterior $q(d_t|h_t, Y_t)$ and Markov prior $p(d_t|d_{t-1})$

Appendix: DAG Acyclicity Constraint (NOTEARS)

Problem: Ensure learned adjacency matrix A represents a DAG (no directed cycles).

Key Insight: Matrix exponential counts paths of all lengths.

$$e^A = I + A + \frac{A^2}{2!} + \frac{A^3}{3!} + \dots$$

- $[A^k]_{ij}$ counts paths of length k from i to j
- Self-loops (cycles): $[A^k]_{ii} > 0$ if node i is on a cycle of length $\leq k$
- Trace: $\text{tr}(A^k) = \sum_i [A^k]_{ii}$ sums all self-loops

NOTEARS Constraint:

$$h(A) = \text{tr}(e^{A \odot A}) - n = 0$$

- $A \odot A$: Element-wise square (ensures non-negative weights)
- If A is a DAG: $\text{tr}(e^{A \odot A}) = n$ (only identity contributes)
- If A has cycles: $\text{tr}(e^{A \odot A}) > n$ (cycle terms contribute)

Optimization: Augmented Lagrangian: $\mathcal{L} + \lambda h(A) + \frac{\rho}{2} h(A)^2$

Appendix: Gumbel-Softmax Trick

Problem: Sampling from categorical distribution is not differentiable.

Step 1: Gumbel-Max Trick (exact sampling)

$$\arg \max_k (\log \pi_k + G_k) \sim \text{Cat}(\pi), \quad \text{where } G_k \sim \text{Gumbel}(0, 1)$$

Step 2: Softmax Relaxation (differentiable approximation)

$$y_k = \frac{\exp((\log \pi_k + G_k)/\tau)}{\sum_j \exp((\log \pi_j + G_j)/\tau)}$$

Temperature τ controls sharpness:

- $\tau \rightarrow 0$: Approaches hard one-hot samples
- $\tau \rightarrow \infty$: Approaches uniform distribution
- Typical: $\tau = 1.0$ during training

Straight-Through Estimator:

- Forward: Use hard samples (argmax)
- Backward: Use soft gradients (softmax)
- Enables discrete structure learning

In CaRS: Used for sampling DAG edges $A_{ij} \in \{0, 1\}$ in differentiable manner.